

This table shows all results in the report. Use the column headers to sort the results in this report. Double-click a result to see detailed metrics. Double-click on demangled names to rename it.

ID	Estimated Speedup [%]	Function Name	Demangled Name	Duration [us] (1477.73 us)	Runtime Improvement [us] (876.07 us)	Compute Throughput [%]	Memory Throughput [%]	# Reg
5	46.55	thresholding_kernel	thresholding_kernel..	144.13	67.09	67.77	67.77	
6	60.64	hough_kernel	hough_kernel..char..	249.70	151.41	42.83	65.92	
7	60.65	get_lines_kernel	get_lines_kernel..	17.92	10.87	17.60	60.90	
8	95.83	draw_lines_kernel	draw_lines_kernel..	168.16	161.15	0.68	0.61	

The following performance optimization opportunities were discovered for this result. Follow the rule links to see more context on the Details page.
Note: Speedup estimates provide upper bounds for the optimization potential of a kernel assuming its overall algorithmic structure is kept unchanged.

- [L1TEX Global Load Access Pattern](#)
Est. Speedup: 60.64%

The memory access pattern for global loads from L1TEX might not be optimal. On average, only 5.3 of the 32 bytes transmitted per sector are utilized by each thread. This could possibly be caused by a stride between threads. Check the [Source Counters](#) section for uncoalesced global loads.
- [Uncoalesced Global Accesses](#)
Est. Speedup: 39.95%

This kernel has uncoalesced global accesses resulting in a total of 1238276 excessive sectors (41% of the total 2987112 sectors). Check the L2 Theoretical Sectors Global Excessive table for the primary source locations. The [CUDA Programming Guide](#) has additional information on reducing uncoalesced device memory accesses.
- [Achieved Occupancy](#)
Est. Speedup: 34.08%

The difference between calculated theoretical (100.0%) and measured achieved occupancy (64.6%) can be the result of warp scheduling overheads or workload imbalances during the kernel execution. Load imbalances can occur between warps within a block as well as across blocks of the same kernel. See the [CUDA Best Practices Guide](#) for more details on optimizing occupancy.